Strategic Features for General Games

Cameron Browne, Dennis J. N. J. Soemers, and Eric Piette
Department of Data Science and Knowledge Engineering (DKE)
Maastricht University, Bouillonstraat 8-10
Maastricht, 6211 LH, The Netherlands

Abstract

This short paper describes an ongoing research project that requires the automated self-play learning and evaluation of a large number of board games in digital form. We describe the approach we are taking to determine relevant features, for biasing MCTS playouts for arbitrary games played on arbitrary geometries. Benefits of our approach include efficient implementation, the potential to transfer learnt knowledge to new contexts, and the potential to explain strategic knowledge embedded in features in human-comprehensible terms.

Introduction

The Digital Ludeme Project[1] is a five-year research project, recently launched at Maastricht University, that aims to model the world's traditional strategy games in a single, playable digital database. This database will be used to find relationships between games and their component *ludemes*,[2] in order to develop a model of the evolution of games throughout recorded human history and to chart their spread across cultures worldwide. This paper describes a new approach to defining strategic features for general games that we are implementing in order to achieve the project's aims.

Statement of the Problem

While there exists substantial archaeological evidence for *ancient games* (i.e. those played before 500AD) the rule sets explaining how these games were actually played are typically lost; our modern understanding of ancient and early games is based primarily on modern reconstructions that have proved to be less than reliable (Murray 1952).

A key challenge in this project will be to evaluate reconstructions of ancient rule sets and to improve them where possible. The general game system we are developing for this project, called LUDII, must therefore be able to:

1. Model the full range of traditional strategy games.
2. Play them at a competent human level.
3. Evaluate rule set reconstructions for: a) historical authenticity, and b) quality of play.
4. Improve rule set reconstructions where needed.

---

[2] Units of game-related information (Borvo 1977).

We aim to model the 1,000 most influential traditional games throughout history, each of which may have multiple interpretations, each of which may require hundreds of variant rule sets to be tested. This is therefore not just a mathematical/computational challenge but also a logistical one, as LUDII may need to be able to learn to play and evaluate up to a million candidate rule sets over the project's five years.

This paper outlines work-in-progress for achieving this ambitious goal, through the self-play learning of simple lightweight features that embody strategic knowledge about the games being modelled. The features we propose are geometry-independent, facilitating their transfer to new contexts, and represent basic strategies that may be understood and explained in human-comprehensible terms.

Geometric Piece Features

The LUDII general game system employs *Monte Carlo tree search* (MCTS) (Browne et al. 2012) as its core method for AI move planning, which has proven to be a superior approach for general games in the absence of domain-specific knowledge (Finnsson and Björnsson 2010). While MCTS performance can vary significantly if strictly random playouts are used (Browne 2013), the reliability of Monte Carlo simulations can be improved by biasing playouts with domain-relevant features.

There is a significant amount of prior research that describes the use of local features or patterns for game-playing agents tailored towards a specific game. Some of the most common approaches for generating features are:

- Manually constructing features using expert knowledge. This was, for example, done by Gelly et al. (2006) for the game of Go, and by Sturtevant and White (2007) for the game of Hearts.

- Exhaustively generating all possible patterns up to a certain restriction (often limited to, e.g., areas of 3×3 cells). This was, for example, done by Gelly and Silver (2007) for the game of Go, and Lorentz and Zosa IV (2017) for the game of Breakthrough. Sturtevant and White (2007) also exhaustively generated combinations (for instance pairs, triples, etc.) of the manually selected features for the game of Hearts as described above.

- Supervised learning based on databases of human expert games. This has been particularly common in the game

of Go (Enderton 1991; Stoutamire 1991; van der Werf et al. 2003; Bouzy and Chaslot 2005; Stern, Herbrich, and Graepel 2006; Araki et al. 2007; Coulom 2007).

Other approaches include Gradual Focus (Skowronski, Björnsson, and Winands 2009), a genetic programming approach to find useful patterns (Hoock and Teytaud 2010), and the use of deep convolutional neural networks for learning relevant features through self-play (Silver et al. 2016).


Figure 1: Bridge completion is a beneficial pattern in Hex.

The features we are interested in here are geometric piece patterns of arbitrary topology. For example, Figure 1 shows a pattern of pieces that represents a beneficial move for connection games played on the hexagonal grid. If Black has just played move **1** to intrude into the virtual connection of two White pieces (left), then White should reply immediately with move **2** to complete that connection. Biasing MCTS playouts to play such moves with higher probability was found to dramatically improve AI playing strength for connection games such as Hex and Y (Raiko and Peltonen 2008).

Feature Definition

Each feature x defines:

1. A *pattern* p_x which specifies one or more elements that must be present or absent in certain locations in game states for feature x to hold.

2. An action a_x which may be encouraged or discouraged in game states where x holds.

For most board games, actions can be sufficiently specified using one or two integers; a location to move "to", and in some games a location to move "from". For example, in games such as Hex and Go a single board location is sufficient to uniquely identify any action in any particular game state. In games such as Draughts or Chess, the current location of the piece to be moved is also required.

Syntax

Each pattern consists of a set of *elements*. We define the following element types for each potential feature element, relative to the current mover:

- "-" Off-board location (for locating edges and corners).
- "." Empty board location.
- "o" Friendly piece (i.e. of the mover's colour).
- "x" Enemy piece (i.e. not of the mover's colour).
- "Pn" Piece belonging to player n.

- "In" Piece with index n in the game definition.[3]

We also define "!" (i.e. "not") versions of these element types. For example, "!P2" means that the specified location does not contain a piece belonging to Player 2. Locations can simultaneously have multiple qualifiers, e.g. a location with both "x" and "!I3" qualified would match an enemy piece at that location unless it has index 3 in the game definition.

Topology

Rather than tying feature topology to any particular board or grid type, feature elements are described by their *relative locations* on the underlying graph of the game board (or more accurately its dual, which defines orthogonal cell adjacencies). Starting with a given *anchor* location and direction, the relative location of each element is defined as a *walk* through the underlying board graph. We assume that the orthogonal neighbours of each board location are listed in consecutive clockwise order, including null placeholders for off-board steps.

Each walk is defined as a sequence of adjacent *steps* through the underlying board graph, where:

- 0 denotes a step forwards in the current direction.
- $-\frac{1}{3}$ denotes $\frac{a}{3}$ counterclockwise turns in a cell with a sides (i.e. one turn in a triangle), then a step forwards,
- $+\frac{2}{4}$ denotes $\frac{2a}{4}$ clockwise turns in a cell with a sides (i.e. two turns in a square), then a step forwards, etc.

Figure 2 shows the relative steps through cells with different numbers of sides. Fractional turns are rounded to the nearest sensible fraction for any given cell (e.g., a turn of $\frac{1}{4}$ in a triangle becomes equal to a turn of $\frac{1}{3}$). The basic mechanism is similar in principle to the use of *turtle commands* in computer graphics.

Note that cells with an odd number of sides, such as the triangular and pentagonal cells shown in Figure 2, do not have a forwards (0) direction. Such cases may be handled by instantiating two pattern instances per ambiguous step: one with 0 replaced by $-\frac{1}{a}$ and one with 0 replaced by $+\frac{1}{a}$ for that step in a cell with a sides.


Figure 2: Steps through various cells based on CW turns.

For example, a knight move may be described using this notation as $\{0, 0, \frac{1}{4}\}$, as shown in Figure 3 (left). This description can be mapped directly to other grids of arbitrary topology to produce plausible results, such as the semi-regular 3.4.6.4 tiling (right). Note that the turn of $\frac{1}{4}$ is ambiguous in the semi-regular grid, and is therefore split into two possible patterns.

[3]The game definition maintains a unique index for each type of

Figure 3: Relative knight move $\{0, 0, \frac{1}{4}\}$ on the square grid (left) and two equivalent moves on the semi-regular 3.4.6.4 grid (right).

Local vs Global Patterns

Patterns may be defined as either:

- *Relative:* Apply to all valid anchor locations across the board (taking null locations into account).
- *Absolute:* Apply only to the specified anchor location.

For convenience, the number of valid rotations and reflections for each pattern may also be specified, so that all possible instances of a given pattern may be generated across the board from a single description.

Reactive vs Proactive Patterns

We distinguish between *reactive* patterns that trigger actions directly in response to the previous player's last move,[4] and *proactive* patterns that can trigger actions anywhere across the board regardless of the previous player's last move. This distinction is similar to the distinction between "responsive" and "non-responsive" patterns used in (Silver et al. 2016).

For example, the bottom row of Figure 4 shows the Hex pattern described earlier in Figure 1 (left) and a reactive pattern that describes this situation (middle). The top row shows a proactive pattern that encourages Hex players to play two disjoint steps away from existing friendly pieces. For clarity in the diagrams, friendly pieces are denoted as white disks and enemy pieces as black discs (the last move made is dotted). White dots indicate empty cells, edges between elements indicate cell adjacency, and green "+" symbols indicate the preferred action triggered by each pattern. The rightmost column of Figure 4 shows how these patterns map easily from the hexagonal grid to the square grid.

Implementation

The LUDII system and associated feature mechanisms are implemented in Java 8. Internal game states are defined using a custom `BitSet` class, called a `ChunkSet`, that compresses the required state information in the minimum number of bits, based on each game's definition.

---

equipment it involves, including *containers* (boards, piles, hands, decks, etc.) and *components* (pieces, tiles, dice, cards, etc).

[4]There may be more than one opponent.

Figure 4: Reactive (bottom) and proactive (top) features.

For example, if a game involves no more equipment than a board and uniform pieces in N colours, then the game state is described by a `ChunkSet` subdivided into *chunks* of B bits per board cell, where B is the lowest power of 2 that provides enough bits to represent every possible state per cell (including state 0 for the empty cells).[5]

For each game, we seek to derive an optimal feature set F, which is the set of features that outperform all other known feature sets for that game on average, within a given computational budget.

Feature Instantiation

As each game is loaded, an *instance* of every possible valid rotation, reflection and translation is pre-generated once for all component features f_x of that game's optimal feature set F, including combinations of valid element types within each feature. Each feature definition can therefore generate hundreds of instances.

Each instance is defined using the same custom `BitSet` class as the game state. Hence, each feature instance can be matched to the given game state efficiently using bitwise parallel operations, such that only a few bitwise operations need to be applied per instance match test, regardless of the feature's complexity.

Feature Application

Feature sets are applied to bias MCTS playouts as follows:

1. For each player move, all legal moves are generated and initialised with equal probability of being selected.

2. Each reactive feature instance corresponding to the previous player's last move is then checked for a match to the current game state, and if a match is detected then

---

[5]Chunk sizes are set to the lowest power of 2 to avoid issues with chunks straddling consecutive `long` member variables.

the feature's weight is added to the probability of the feature's action being selected (note that feature weights can be negative for detrimental moves).

3. Each proactive feature instance is then applied across the board and checked for a match to the current game state, and if a match is detected then the feature's weight is added to the probability of the feature's action being selected.

4. The move to be made is then selected randomly according to the resulting biased distribution.

Features can similarly be used in other parts of MCTS, such as its selection phase. Note that reactive features are much more efficient to apply than proactive ones, as only those patterns relevant to the previous player's last move need be tested. We therefore seek to generate reactive features, and to replace proactive features with their reactive equivalents, where possible.

Feature Generation

Manually generating features using expert knowledge, or extracting them through supervised learning from human expert games, are not suitable solutions for the current task in which a wide variety of unfamiliar games must be supported. Exhaustive generation of features can be done more easily, but has clear drawbacks due to the massive volume of features it can generate.

One advantage of describing games by their component ludemes is that these may be exploited to generate a plausible set of candidate feature patterns based on the specified rules and equipment. Such patterns may not be suitable features in and of themselves, but represent "minimum" patterns which must be present in every feature – for example, every Hex or Go feature must contain an empty "to" action location – allowing significant reductions in the number of possible patterns to be generated. For many games, this kind of knowledge can be extracted automatically from ludemes.

The subsequent fine-tuning of features to obtain the optimal feature set F for a given game is still a work-in-progress. Suffice it to say that *frequent pattern mining* (Aggarwal and Han 2014) approaches to iteratively building features from large numbers of randomly generated self-play games has not proved effective. Learning feature sets against a random opponent can have disastrous – and sometimes amusing – results in the unexpected behaviours that can thwart the random player. There does not seem to be any escape from evaluating feature sets using more time-consuming but realistic tests against intelligent AI opponents, with MCTS as the baseline yardstick.

Features and Strategies

An attractive aspect of the proposed approach is that learnt features have the potential to encode simple strategies for the games to which they apply. For example, the features shown on the top row of Figure 5 encourage the player to make lines of four pieces of their colour, while the features shown on the bottom row encourage the player to *not* make lines of three pieces of their colour (the red "–" symbols indicate negative weights that discourage such actions).


Figure 5: Patterns for a "line of 4" strategy (top) and a "not line of 3" strategy (bottom).

Figure 6: Patterns for a "form groups of 3" strategy.

The features shown in Figure 6 encourage the player to form groups of three pieces of their colour, by encouraging singleton pieces to grow in any direction then discouraging growth beyond group size 3. Similarly, the features shown in Figure 7 encourage the player to form long, thin groups of their pieces, by encouraging singleton pieces to grow in any direction, then encouraging friendly pairs to extend at the ends but discouraging growth at the common points adjacent to both pieces (which would create shorter, thicker groups).

Feature Explanation

An additional benefit of the ludemic model for games is that features applicable to a given game may be readily transferred as candidate features for related games defined by similar ludemes. Any feature that has proved relevant for a given context is a good starting point for similar contexts.

Further, we can potentially invert the causal relationship between ludemes and features to reverse engineer comprehensible explanations for the learnt strategies that optimal feature sets represent, by determining *which* ludemes are responsible for given features. This relationship between a game's ludemes and its derived features is further strengthened by the fact that initial candidate feature sets are derived directly from the game's ludemic description during feature generation process.

Each ludeme in the LUDII Ludeme Library corresponds to a Java class with an appropriate name, providing a source of convenient plain English captions for learnt concepts. And even if the ludeme names do not provide sufficient expla-


Figure 7: Patterns for a "form long thin groups" strategy.

nation in themselves, they provide hints for further post-processing steps to find geometric relationships within the feature patterns. It is plausible that simple descriptions such as "make lines of 4" or "make groups of 3" may be derived from such feature sets. Implementing such mechanisms for explainable AI in the context of startegy learning for general games will be a key focus of future work.

Conclusion

Lightweight features based on geometric piece patterns have a number of advantages for our work on the Anonymous Project. They improve MCTS playing strength, map readily to other geometries, encode simple strategies, can be associated with the underlying ludemic descriptions of games, and have the potential to help explain learnt strategies in human-comprehensible terms. We will continue developing and testing this approach as the LUDII general game system matures to support an increasing range of game types, and investigating the automated extraction and explanation of relevant strategies from such learnt features.

Acknowledgements

References

Aggarwal, C. C., and Han, J., eds. 2014. *Frequent Pattern Mining*. Springer.

Araki, N.; Yoshida, K.; Tsuruoka, Y.; and Tsujii, J. 2007. Move prediction in Go with the maximum entropy method. In *Proceedings of the 2007 IEEE Symposium on Computational intelligence and Games*, 189–195. IEEE.

Borvo, A. 1977. *Anatomie d'un jeu de cartes: L'Aluette ou le Jeu de Vache*. Nantes: Librairie Nantaise Yves Vachon.

Bouzy, B., and Chaslot, G. 2005. Bayesian generation and integration of K-nearest-neighbor patterns for 19x19 Go. In Kendall, G., and Lucas, S., eds., *Proceedings of the 2005 IEEE Symposium on Computational Intelligence in Games*, 176–181. IEEE.

Browne, C.; Powley, E.; Whitehouse, D.; Lucas, S.; Cowling, P. I.; Rohlfshagen, P.; Tavener, S.; Perez, D.; Samothrakis, S.; and Colton, S. 2012. A survey of Monte Carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in Games* 4(1):1–49.

Browne, C. 2013. A problem case for UCT. *IEEE Transactions on Computational Intelligence and AI in Games* 5(1):69–74.

Coulom, R. 2007. Computing "ELO ratings" of move patterns in the game of Go. *ICGA Journal* 30(4):198–208.

Enderton, H. D. 1991. The Golem Go program. Technical Report CMU-CS-92-101, School of Computer Science, Carnegie-Mellon University.

Finnsson, H., and Björnsson, Y. 2010. Learning simulation control in general game-playing agents. In *Proceedings of the Twenty-Fourth AAAI Conference on Artificial Intelligence*, 954–959. AAAI Press.

Gelly, S., and Silver, D. 2007. Combining online and offline knowledge in UCT. In *Proceedings of the 24th International Conference on Machine Learning*, 273–280.

Gelly, S.; Wang, Y.; Munos, R.; and Teytaud, O. 2006. Modification of UCT with patterns in Monte-Carlo Go. Technical Report RR-6062, INRIA, Paris.

Hoock, J.-P., and Teytaud, O. 2010. Bandit-based genetic programming. In Esparcia-Alcázar, A. I.; Ekárt, A.; Silva, S.; Dignum, S.; and Uyar, A. Ş., eds., *European Conference on Genetic Programming*, volume 6021 of *Lecture Notes in Computer Science*, 268–277. Springer.

Lorentz, R. J., and Zosa IV, T. E. 2017. Machine learning in the game of Breakthrough. In Winands, M. H. M.; van den Herik, H.; and Kosters, W. A., eds., *Advances in Computer Games*, volume 10664 of *Lecture Notes in Computer Science*, 140–150. Springer.

Murray, H. 1952. *A History of Board-Games Other Than Chess*. Clarendon Press.

Raiko, T., and Peltonen, J. 2008. Application of UCT search to the connection games of Hex, Y, *Star, and Renkula! In *Proceedings of the Finnish Artificial Intelligence Conference*, 89–93.

Silver, D.; Huang, A.; Maddison, C.; Guez, A.; Sifre, L.; van den Driessche, G.; Schrittwieser, J.; Antonoglou, I.; Panneershelvam, V.; Lanctot, M.; Dieleman, S.; Grewe, D.; Nham, J.; Kalchbrenner, N.; Sutskever, I.; Lillicrap, T.; Leach, M.; Kavukcuoglu, K.; Graepel, T.; and Hassabis, D. 2016. Mastering the game of Go with deep neural networks and tree search. *Nature* 529(7587):484–489.

Skowronski, P.; Björnsson, Y.; and Winands, M. H. M. 2009. Automated discovery of search-extension features. In van den Herik, H. J., and Spronck, P., eds., *Advances in Computer Games*, volume 6048 of *Lecture Notes in Computer Science*. Springer, Berlin, Heidelberg.

Stern, D.; Herbrich, R.; and Graepel, T. 2006. Bayesian pattern ranking for move prediction in the game of Go. In Cohen, W. W., and Moore, A., eds., *Proceedings of the 23rd International Conference on Machine Learning*, 873–880.

Stoutamire, D. 1991. Machine learning, game play, and Go. Technical Report TR 91-128, Center for Automation and Intelligent Systems Research, Case Western Reserve University.

Sturtevant, N. R., and White, A. M. 2007. Feature construction for reinforcement learning in Hearts. In van den Herik, H. J.; Ciancarini, P.; and Donkers, H. H. L. M., eds., *Computers and Games*, volume 4630 of *Lecture Notes in Computer Science*, 122–134. Springer.

van der Werf, E.; Uiterwijk, J. W. H. M.; Postma, E.; and van den Herik, J. 2003. Local move prediction in Go. In Schaeffer, J.; Müller, M.; and Björnsson, Y., eds., *Computers and Games*, volume 2883 of *Lecture Notes in Computer Science*. Springer.